



# Introduction to CNNs



# Deep Learning

- Welcome to the section on Convolutional Neural Networks!
- CNNs are a specific architecture of Neural Networks that are extremely effective at dealing with image data.
- Let's review what we will learn in this section!



# Deep Learning

- CNN Section
  - Understanding CNNs
    - Image Kernels and Filters
    - Convolutions
    - Pooling Layers
  - MNIST Dataset (Grayscale Images)
  - CIFAR-10 Dataset (Color Images)



# Deep Learning

- CNN Section
  - Working with image files (jpg,png,etc...) with CNNs.
  - CNNs for malaria blood cell classification.
  - CNN exercise on Fashion Image Dataset.



# MNIST Data



# Deep Learning

- A classic data set in Deep Learning is the MNIST data set.
- Let's quickly cover some basics about it since we'll be using it quite frequently during this section of the course!



## Deep Learning

- Fortunately this data is easy to access with PyTorch Vision. We will download:
  - 60,000 training images
  - 10,000 test images



# Deep Learning

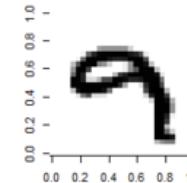
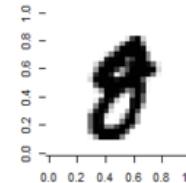
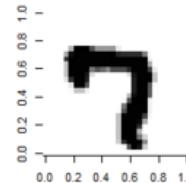
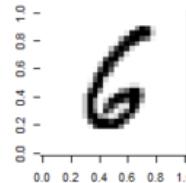
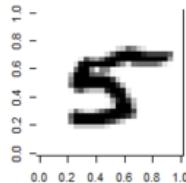
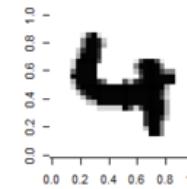
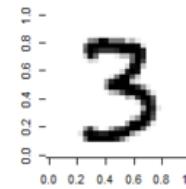
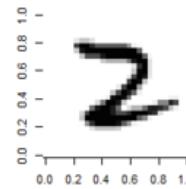
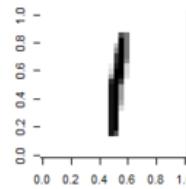
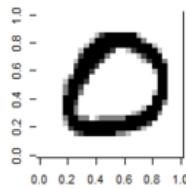
- The MNIST data set contains handwritten single digits from 0 to 9





# Deep Learning

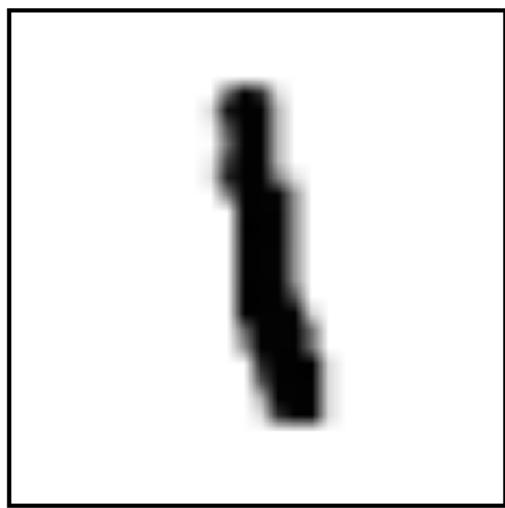
- A single digit image can be represented as an array





# Deep Learning

- Specifically, 28 by 28 pixels



$\approx$

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	.6	.8	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	.7	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	.7	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	.5	1	.4	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	1	.4	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	.7	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	.9	1	.1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	.3	1	.1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0



# Deep Learning

- The values represent the grayscale image

$$\begin{array}{|c|} \hline \text{A grayscale image of a handwritten digit '4' on a white background.} \\ \hline \end{array} \approx \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & .6 & .8 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & .7 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & .7 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & .5 & 1 & .4 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & .4 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & .7 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & .9 & 1 & .1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & .3 & 1 & .1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



# Deep Learning

- We will first explore how we could approach this data set with a standard ANN.
- In this case, to feed it into our network we will need to flatten the 28 by 28 array to 784

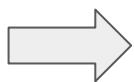


# Deep Learning

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.6	.8	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	.1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	.1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.5	.1	.4	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.4	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.4	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.7	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.1	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.9	.1	.1	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.3	.1	.1	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0



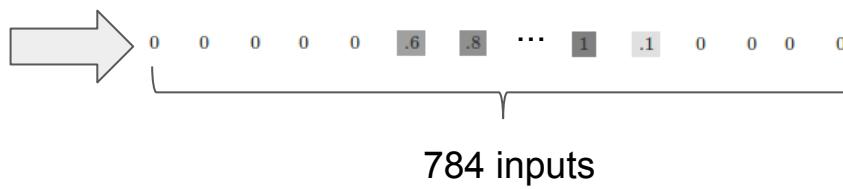
# Deep Learning



0 0 0 0 0 .6 .8 ... 1 .1 0 0 0 0

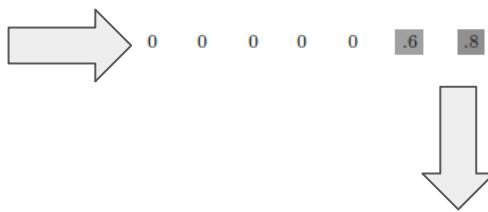


# Deep Learning



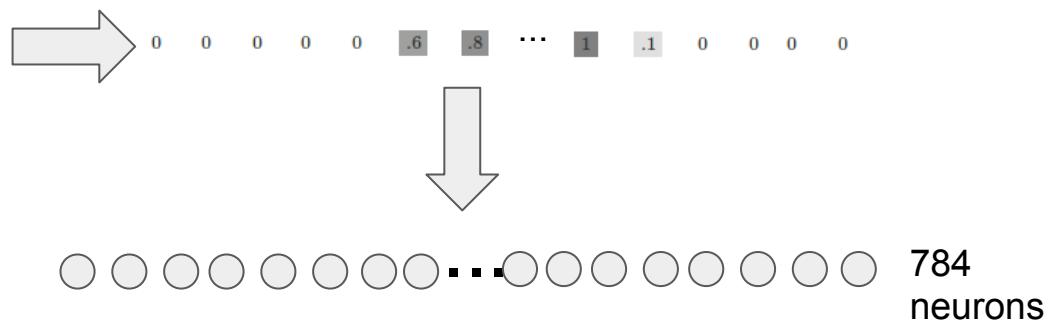


# Deep Learning



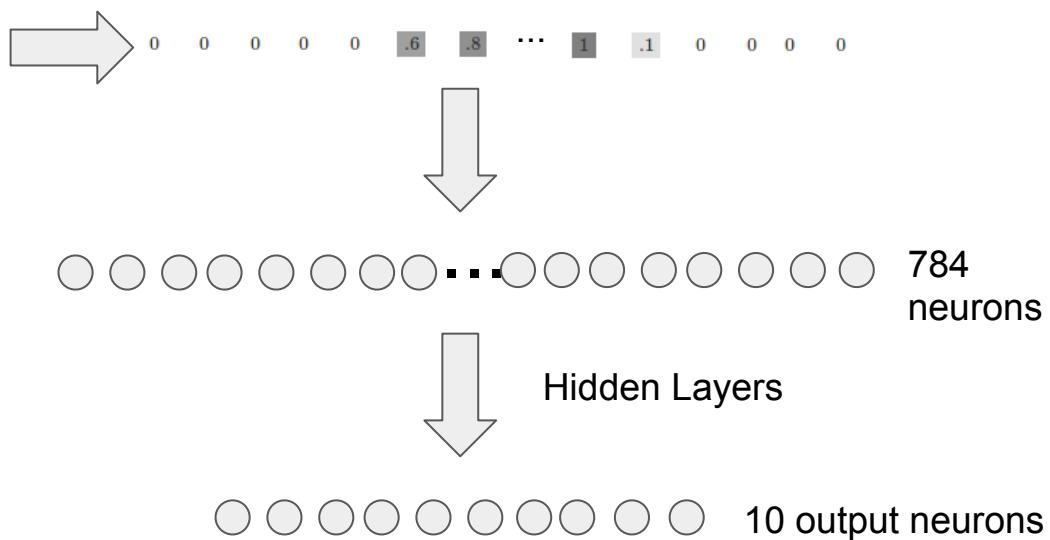


# Deep Learning





# Deep Learning





# Deep Learning

- Flattening out the image ends up removing some of the 2-D information, such as the relationship of a pixel to its neighboring pixels.
- For now, we'll ignore this, but come back to it later when we discuss CNN in depth!



# MNIST ANN

PART ONE: DATA



# MNIST ANN

PART TWO: NEURAL NETWORK



# MNIST ANN

## PART THREE: TRAINING AND EVALUATION



# MNIST ANN

## PART FOUR: EVALUATION



# Image Filters



# Deep Learning

- Before we dive into CNNs, let's first discuss a few key ideas in computer vision.
- Computer vision is a general term of using computer programs to process image data.



# Deep Learning

- If you've ever used photo editing software, you have probably seen filters, such as a blur filter. But how do these work?





# Deep Learning

- Filters are essentially an **image kernel**, which is a small matrix applied to an entire image.
- Certain popular filters are well known, for example a blur filter:

$$\begin{pmatrix} 0.0625 & 0.125 & 0.0625 \\ 0.125 & 0.25 & 0.125 \\ 0.0625 & 0.125 & 0.0625 \end{pmatrix}$$



# Deep Learning

- Let's explore how these image kernel/filters actually get applied to an image.



# Deep Learning

- Filters allow us to transform images

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0



# Deep Learning

- Here we have an grayscale image

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0



# Deep Learning

- Values scaled between -1 and 1

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0



# Deep Learning

- Filters allow us to transform images

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

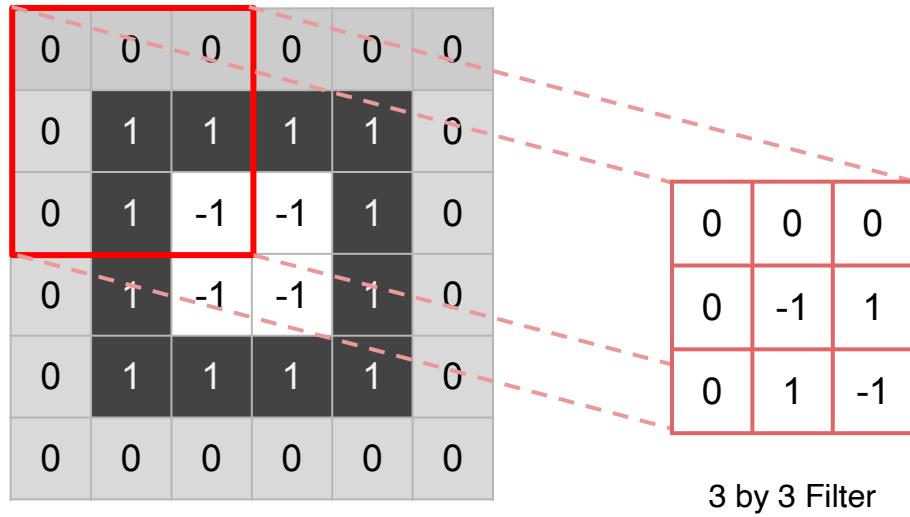
0	0	0
0	-1	1
0	1	-1

3 by 3 Filter



# Deep Learning

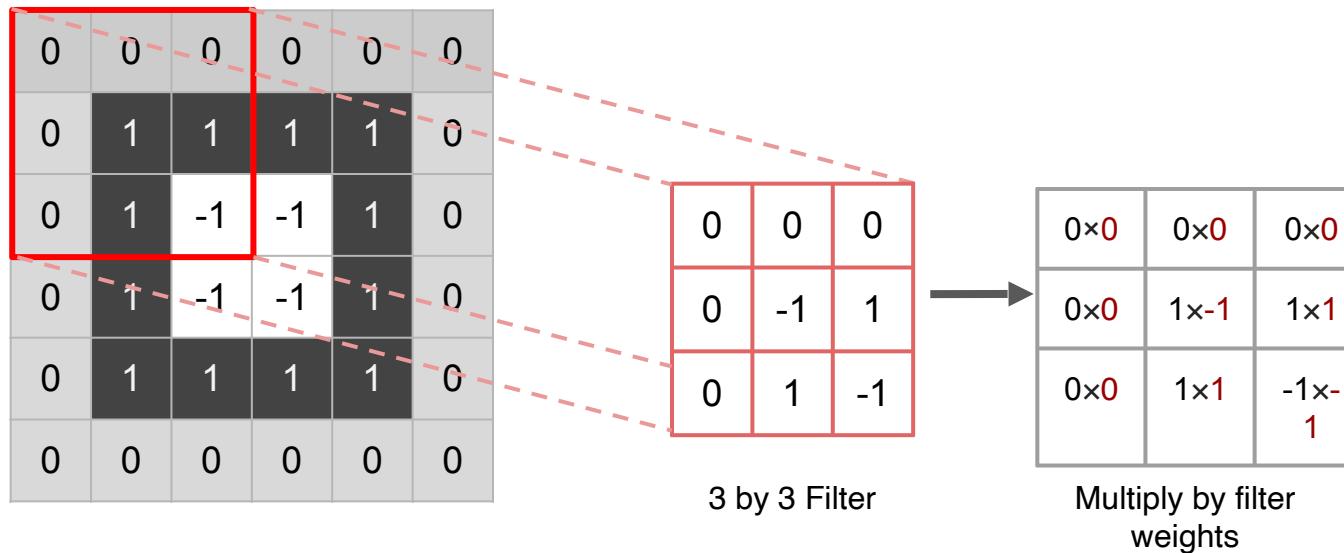
- Filters allow us to transform images





# Deep Learning

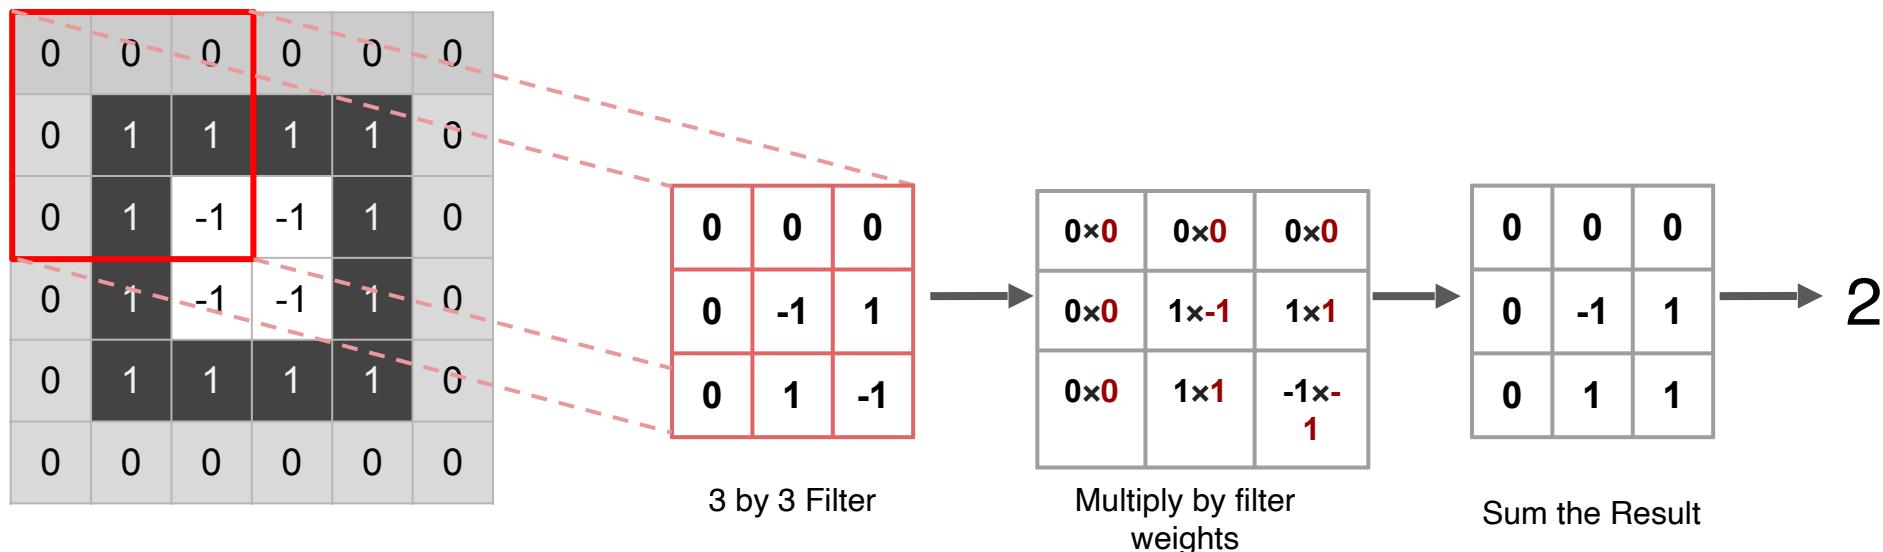
- Filters are essentially a matrix





# Deep Learning

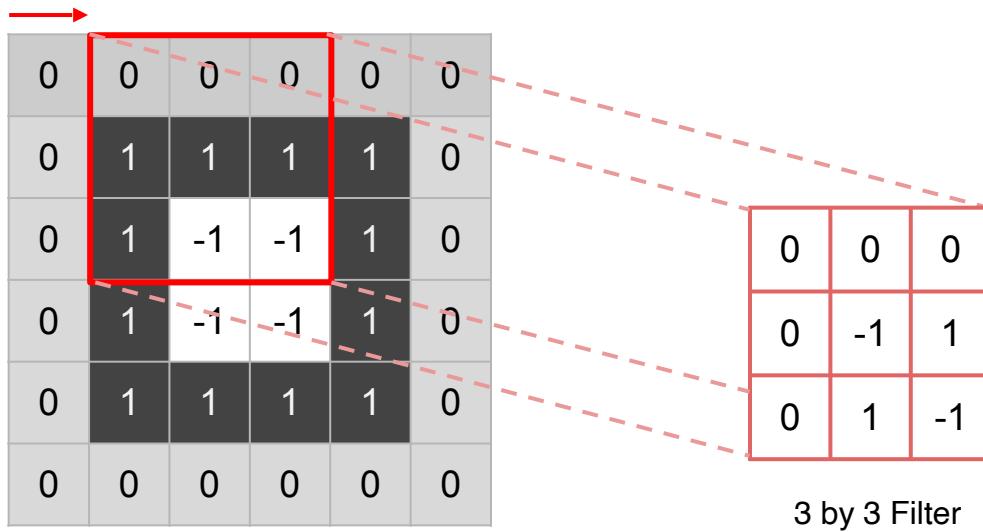
- Notice how the resolution will decrease





# Deep Learning

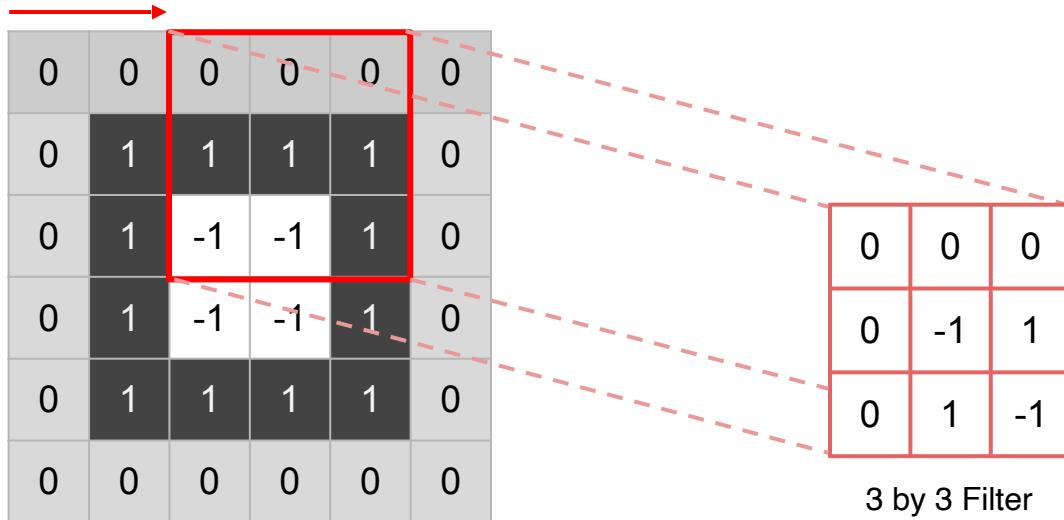
- We can also edit our stride distance





# Deep Learning

- Stride Distance of 2 Example





# Deep Learning

- Let's explore an interactive example!
  - **[setosa.io/ev/image-kernels/](https://setosa.io/ev/image-kernels/)**



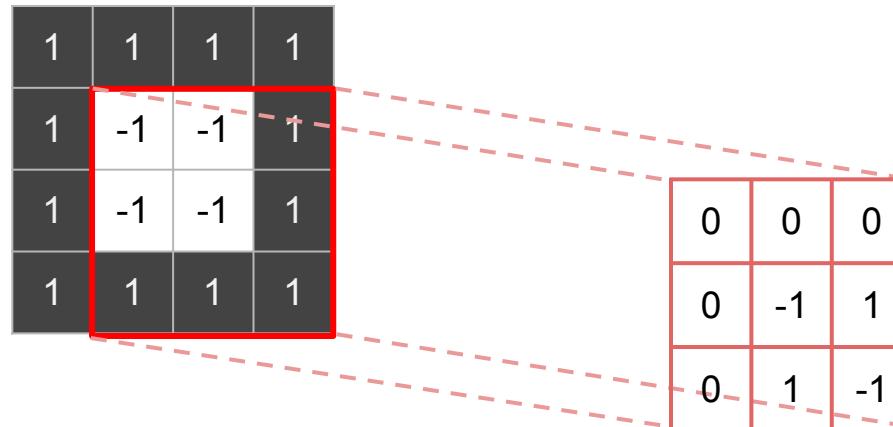
# Deep Learning

- In the context of CNNs, these “filters” are referred to as **convolution kernels**.
- The process of passing them over an image is known as **convolution**.
- Let’s go over a few more important factors!



# Deep Learning

During convolution, we would lose borders

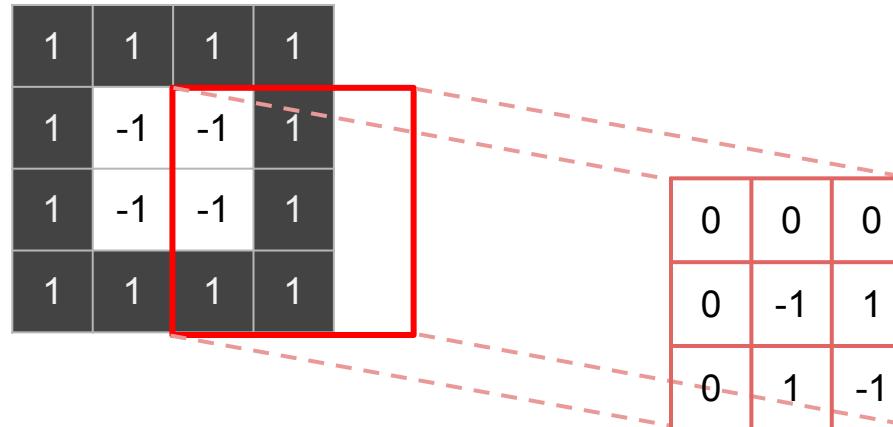


3 by 3 Filter



# Deep Learning

During convolution, we would lose borders

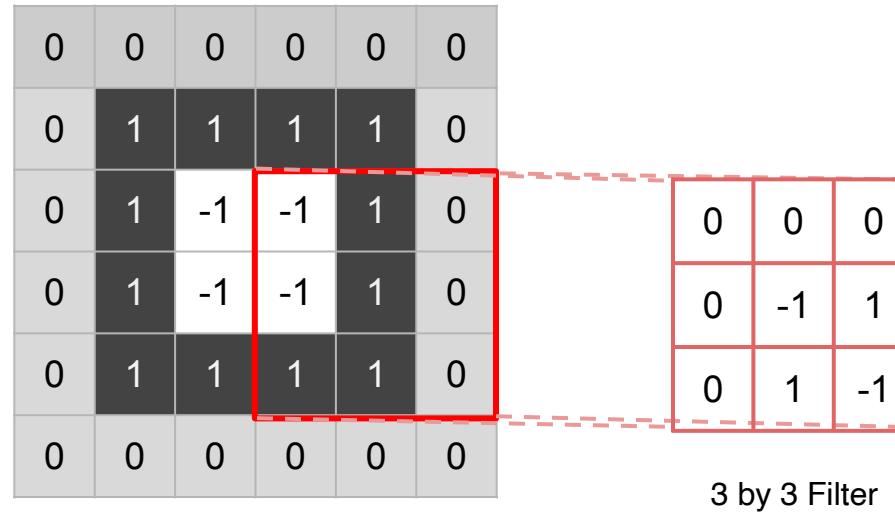


3 by 3 Filter



# Deep Learning

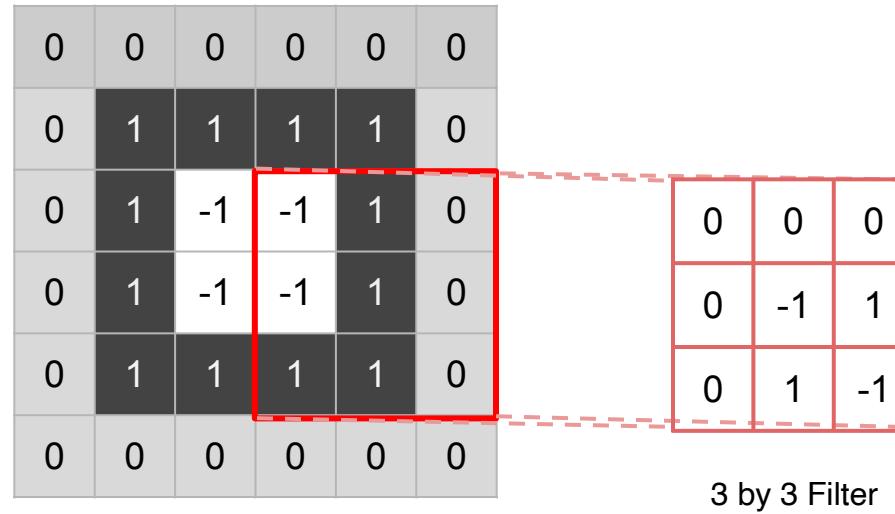
We can **pad** the image with more values





# Deep Learning

This allows us to preserve the image size





# Deep Learning

- Now that we understand image filters, let's explore the architecture of a CNN that allows the network to come up with the best weights for a filter in the **convolutional layer**.



# Convolutional Layers



# Deep Learning

- Recall that running an ANN for the MNIST data set resulted in a network with relatively good accuracy.
- However, there are some issues with always using ANN models for image data.



# Deep Learning

- ANNs
  - Large amount of parameters (over 100,000 for tiny 28 by 28 images)
  - We lose all 2D information by flattening out the image
  - Will only work on very similar, well centered images.



# Deep Learning

- A CNN can use convolutional layers to help alleviate these issues.
- A convolutional layer is created when we apply multiple image filters to the input images.
- The layer will then be trained to figure out the best filter weight values.



# Deep Learning

- A CNN also helps reduce parameters by focusing on **local connectivity**.
- Not all neurons will be fully connected.
- Instead, neurons are only connected to a subset of local neurons in the next layer (these end up being the filters!)



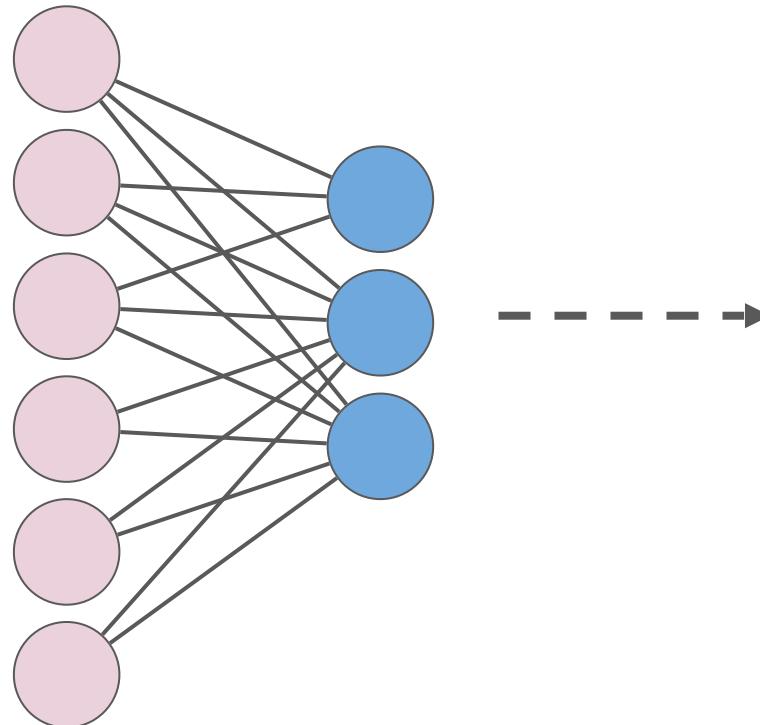
# Deep Learning

- Let's understand this local connectivity and its connection to filters by beginning with a simplified 1-D example.
- We will later expand this to 2-D inputs for a grayscale image and then later on to 3-D tensor inputs for color images.



# Deep Learning

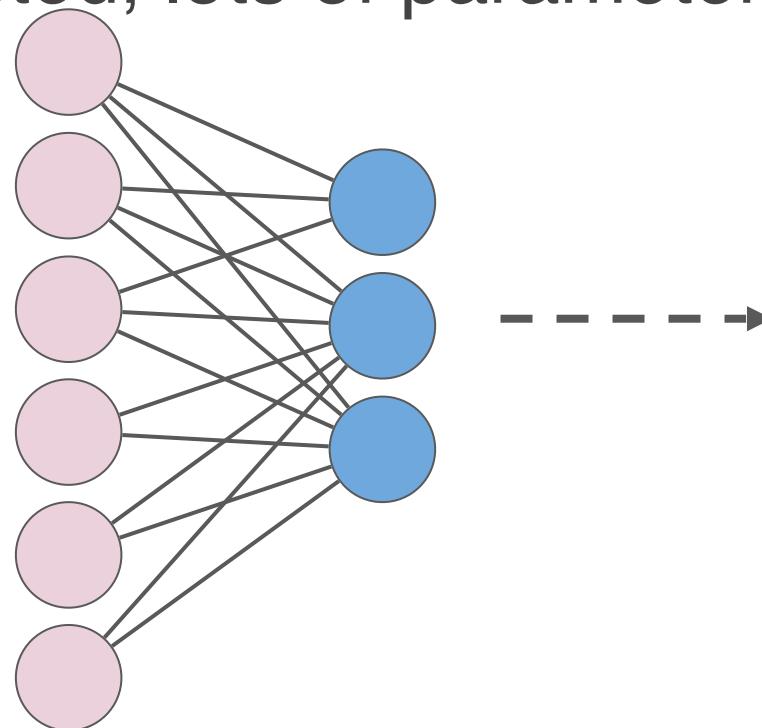
- ANN





# Deep Learning

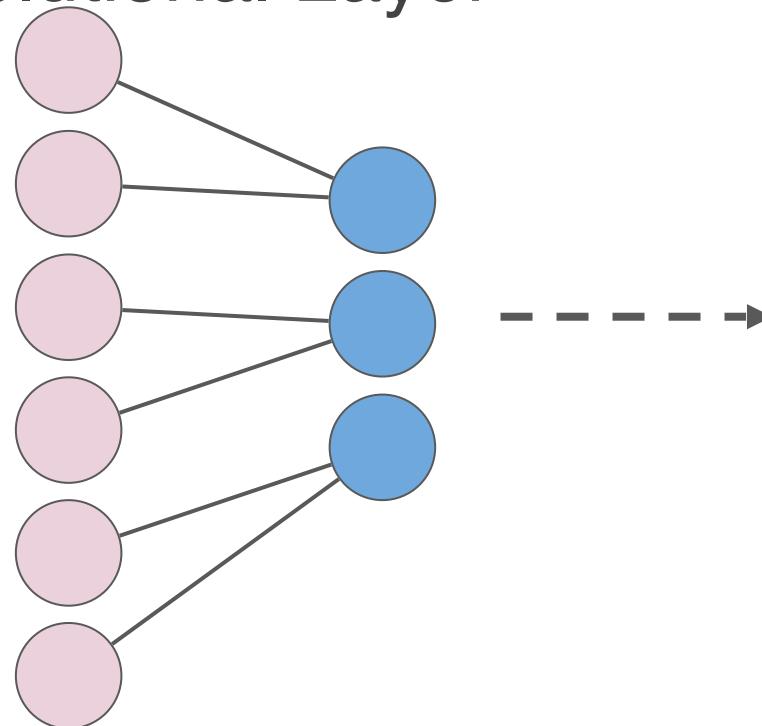
- Fully Connected, lots of parameters!





# Deep Learning

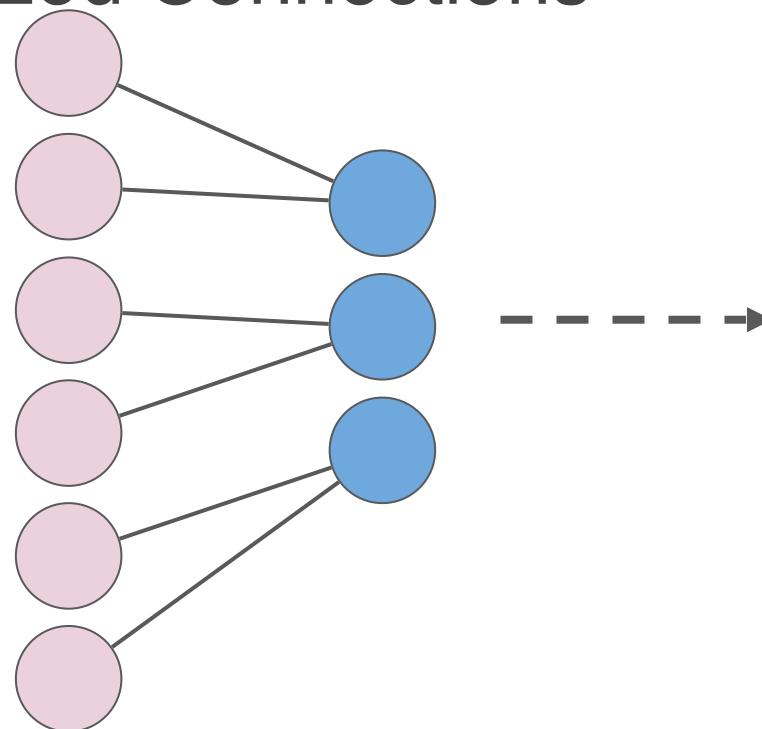
- CNN - Convolutional Layer





# Deep Learning

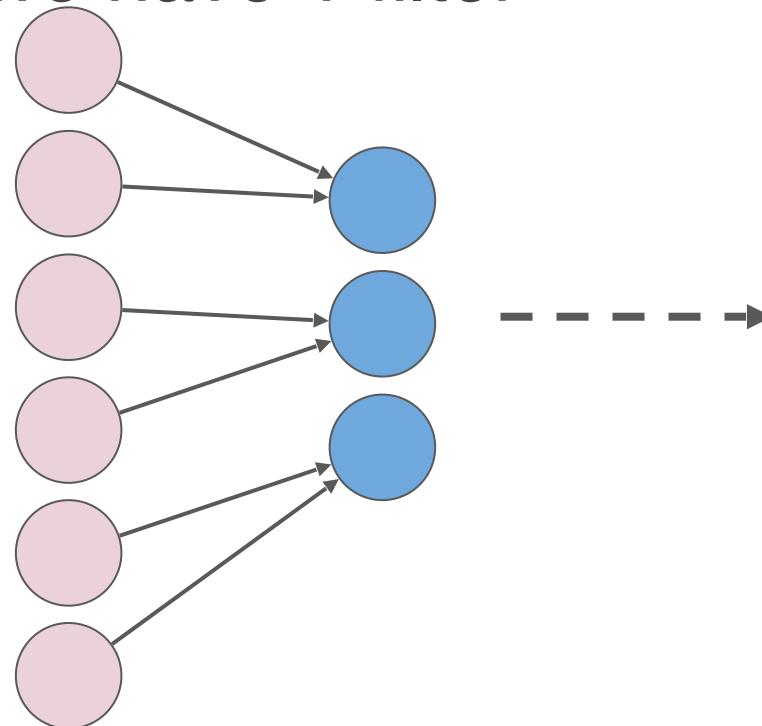
- CNN - Localized Connections





# Deep Learning

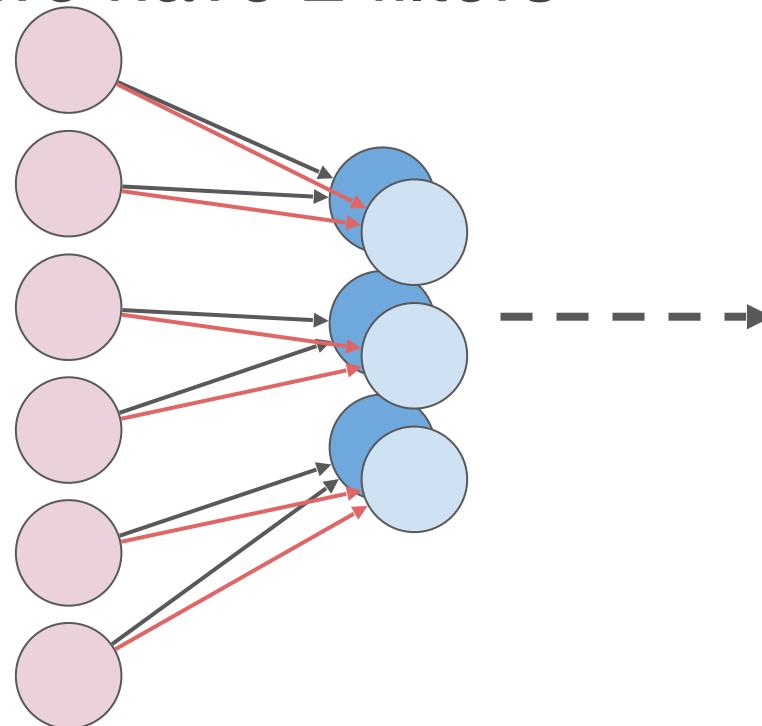
- CNN - Here we have 1 filter





# Deep Learning

- CNN - Here we have 2 filters



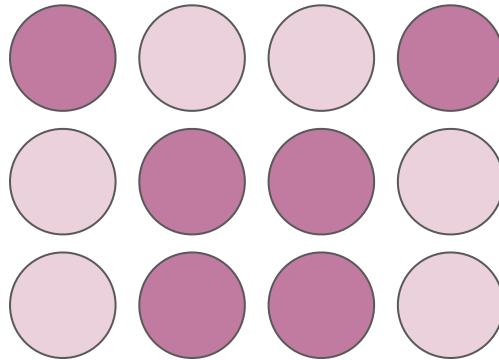


## Deep Learning

- We just saw a 1-D example of convolution, but recall, grayscale images are 2-D!
- Plus we want to preserve that 2-D relational information in the convolutional layer.



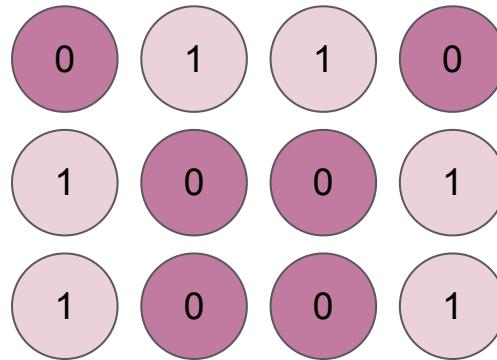
# Deep Learning



**Input Image**



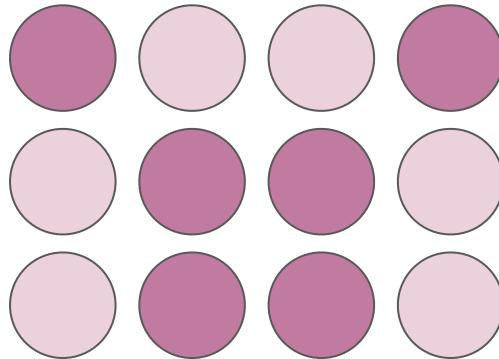
# Deep Learning



**Input Image**



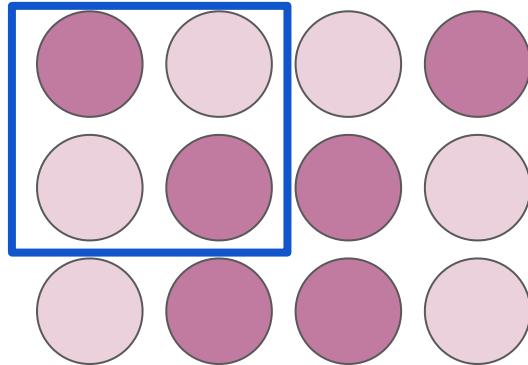
# Deep Learning



**Input Image**



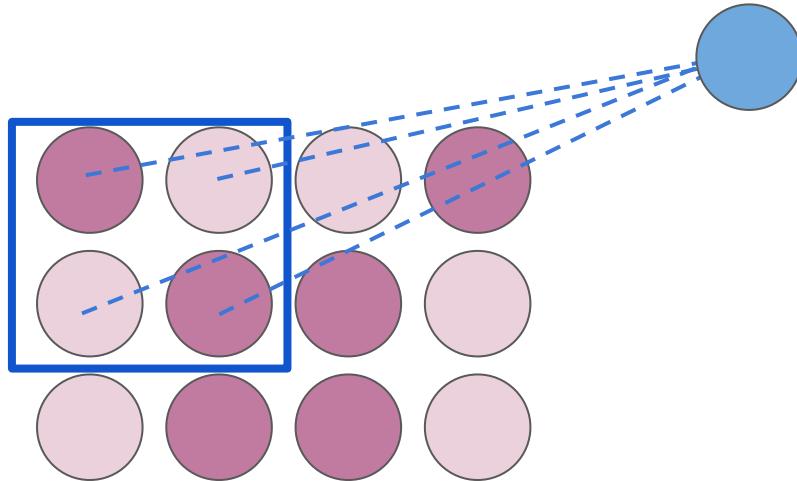
# Deep Learning



**Input Image**



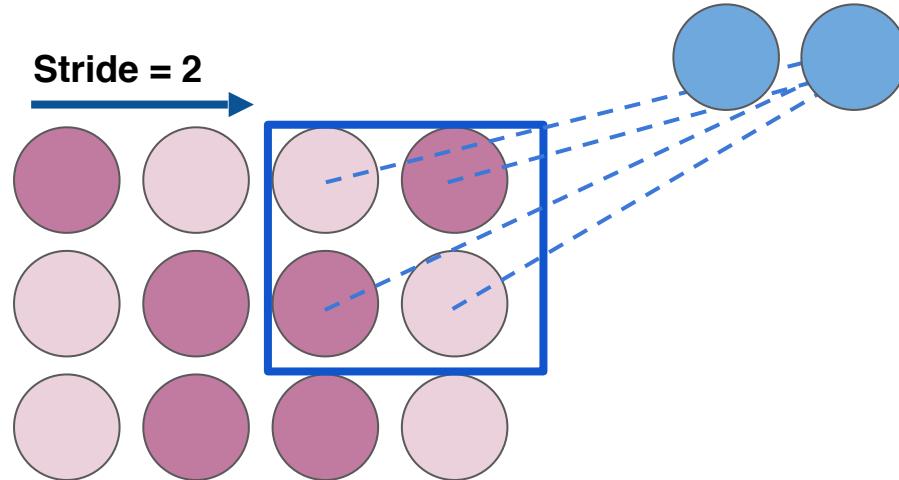
# Deep Learning



**Input Image**



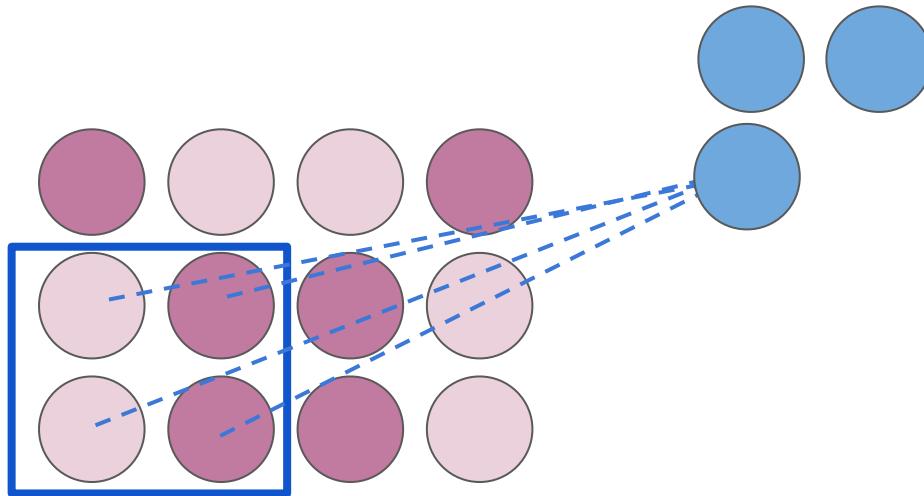
# Deep Learning



**Input Image**



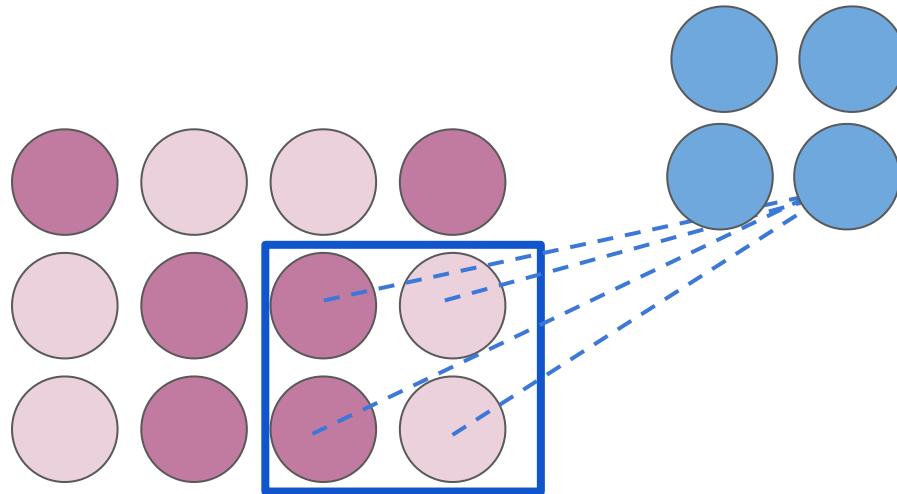
# Deep Learning



**Input Image**



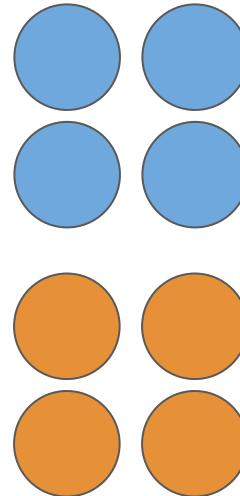
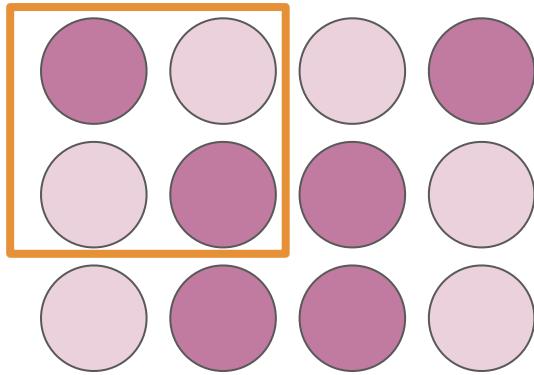
# Deep Learning



**Input Image**



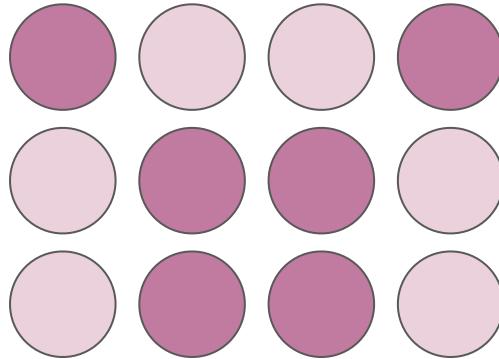
# Deep Learning



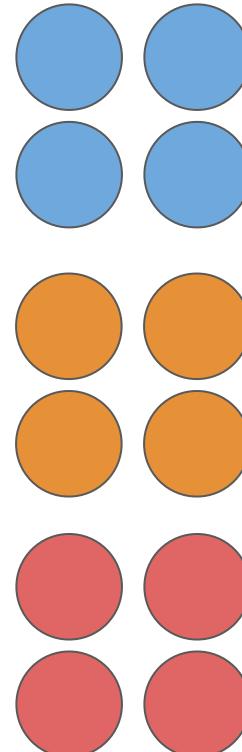
**Input Image**



# Deep Learning

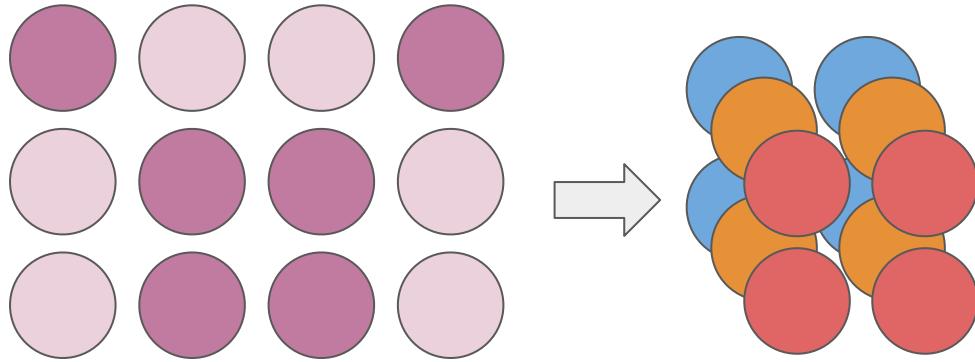


**Input Image**





# Deep Learning



**Input Image**

**Convolutional  
Layer**

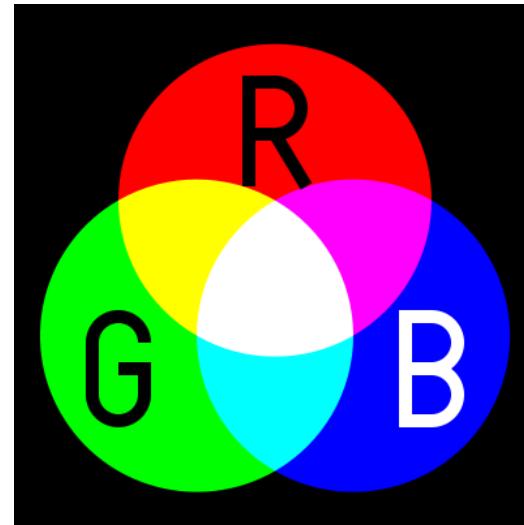


# Deep Learning

- But this was just in 2D for grayscale images.
- What about color images?
- Color Images can be thought of as 3D Tensors consisting of Red, Green, and Blue color channels.

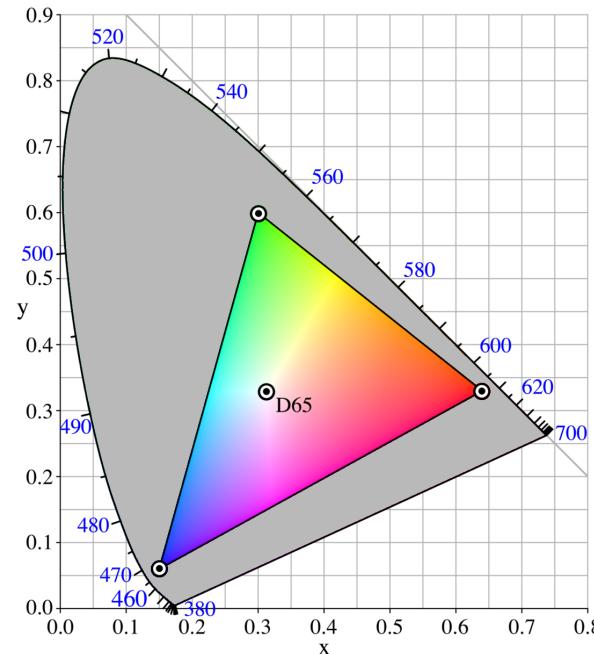


- Additive color mixing allows us to represent a wide variety of colors by simply combining different amounts of R, G, B.



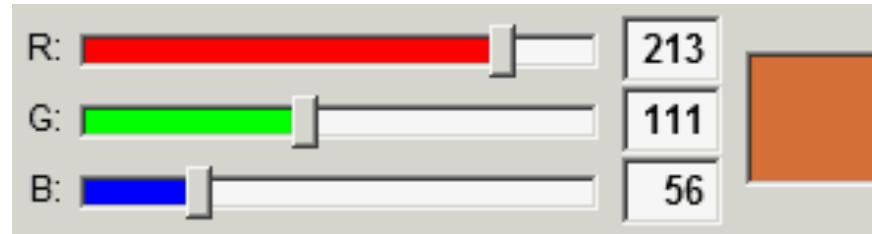


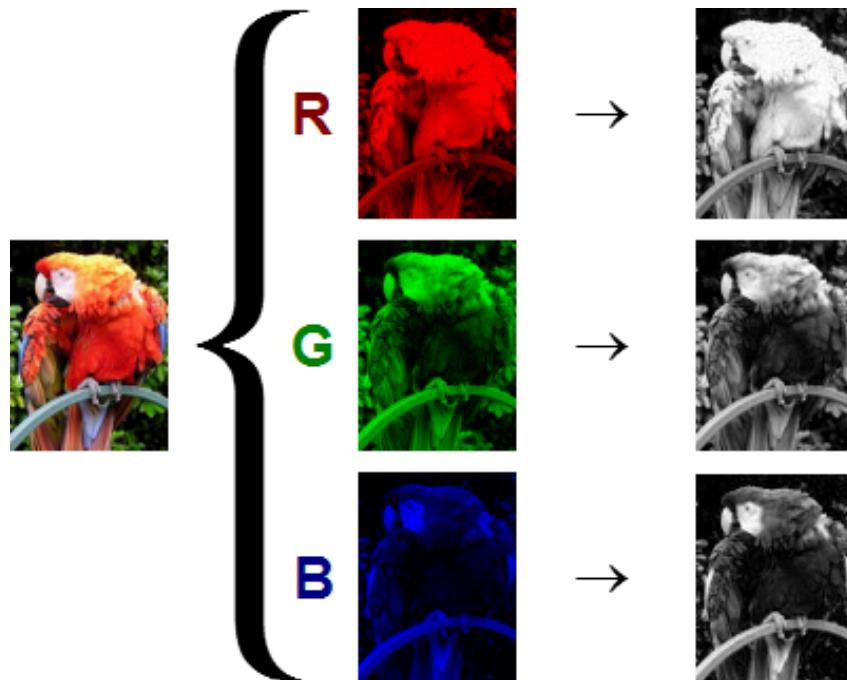
- RGB allows to produce a range of colors





- Each color channel will have intensity values.
- You may have already seen this sort of representation in other software with RGB sliders.





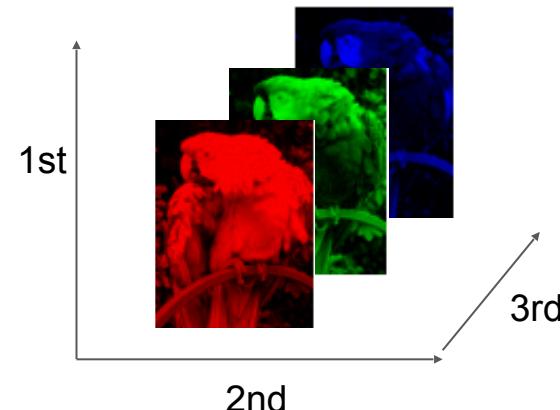
- The shape of the color array then has 3 dimensions.
  - Height
  - Width
  - Color Channels

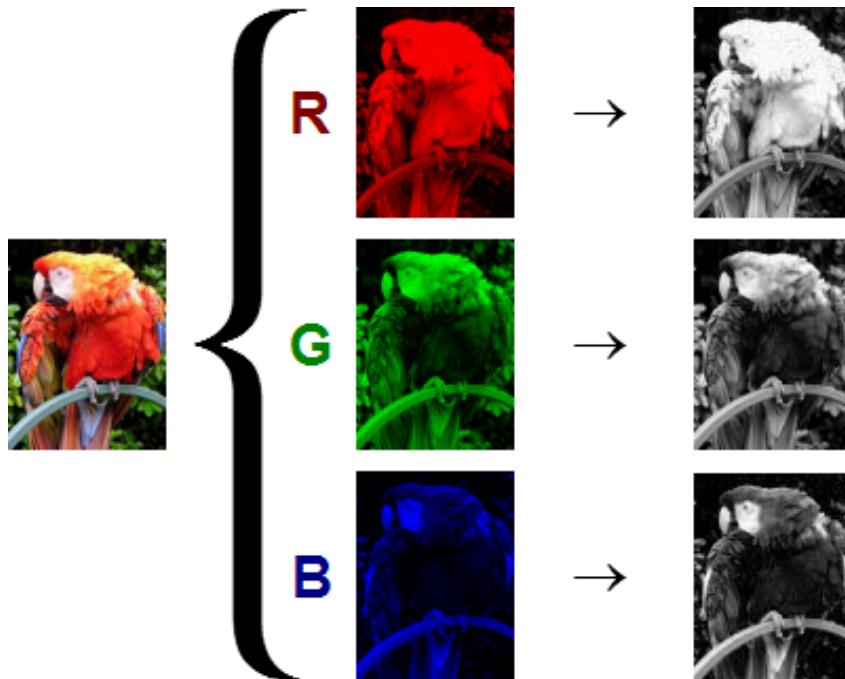


- This means when you read in an image and check its shape, it will look something like:
  - **(1280,720,3)**
  - **1280** pixel width
  - **720** pixel height
  - **3** color channels



- This means when you read in an image and check its shape, it will look something like:
  - **(720,1280,3)**
    - 720 pixel height
    - 1280 pixel width
    - 3 color channels





- Keep in mind the computer won't "know" a channel is Red, it just knows that there are now 3 intensity channels.

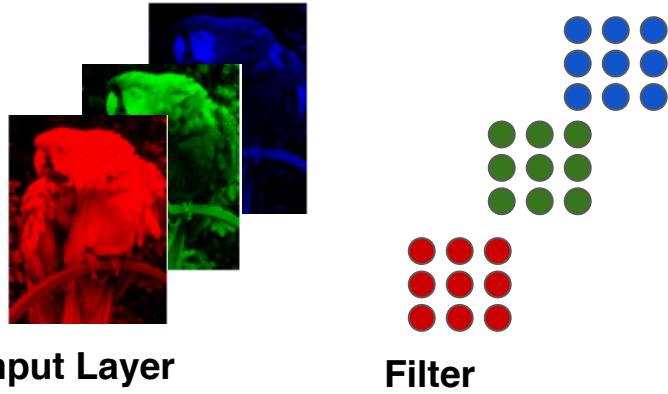


## Deep Learning

- How do we then perform a convolution on a color image?
- We end up with a 3D filter, with values for each color channel.

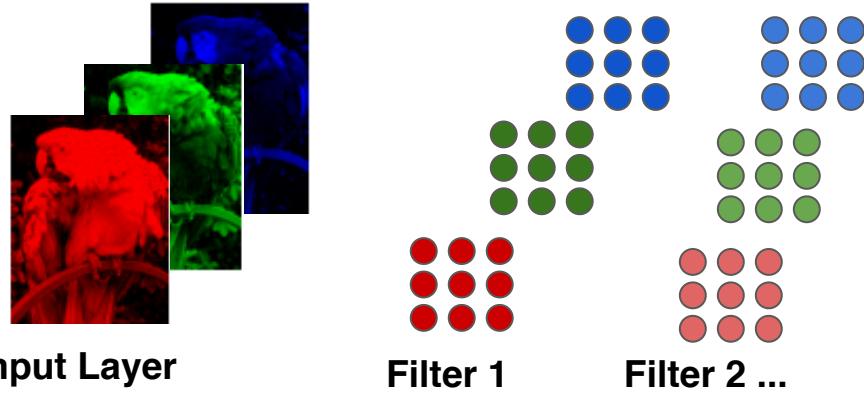


# Deep Learning





# Deep Learning





# Deep Learning

- Often convolutional layers are fed into another convolutional layer.
- This allows the networks to discover patterns within patterns, usually with more complexity for later convolutional layers.



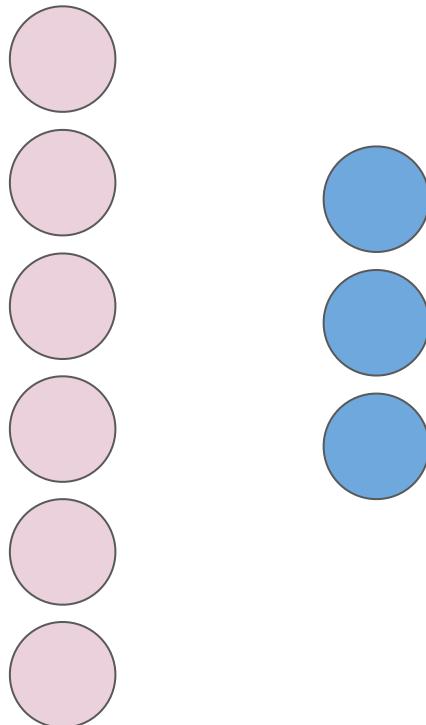
# Deep Learning

- We've learned a lot so far and we have one final theory topic to cover before coding our own CNNs, and that is the Pooling Layer! (also known as a downsampling layer)



# Deep Learning

- ANN





# Pooling Layers



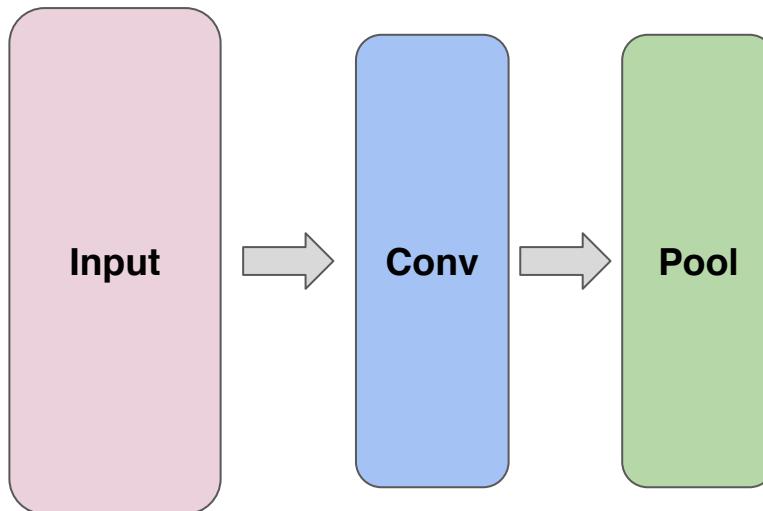
# Deep Learning

- Even with local connectivity, when dealing with color images and possibly 10s or 100s of filters we will have a large amount of parameters.
- We can use pooling layers to reduce this.



# Deep Learning

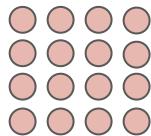
- Pooling layers accept convolutional layers as input:





# Deep Learning

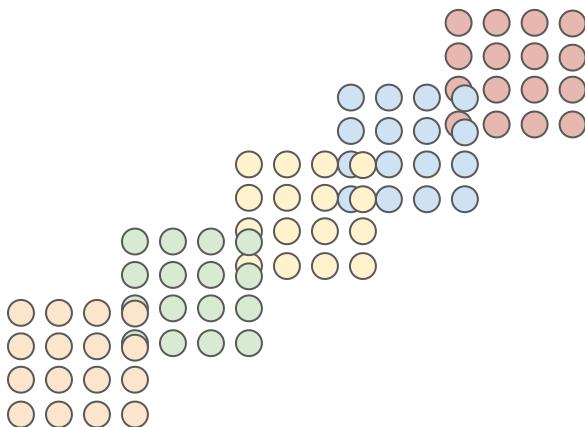
- Our convolutional layers will often have many filters





# Deep Learning

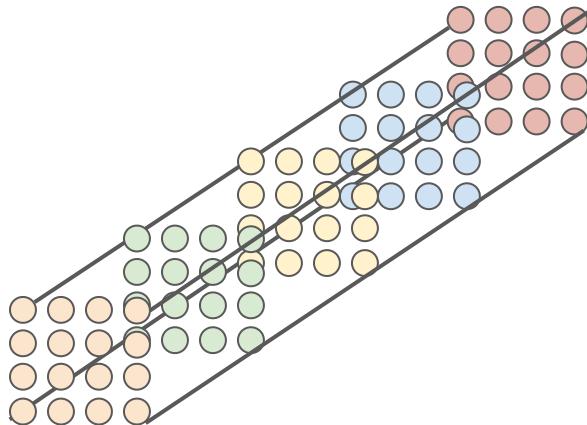
- Our convolutional layers will often have many filters





# Deep Learning

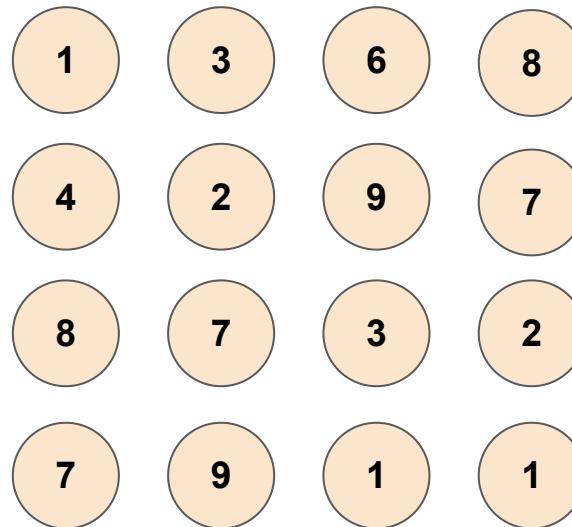
- Our convolutional layers will often have many filters





# Deep Learning

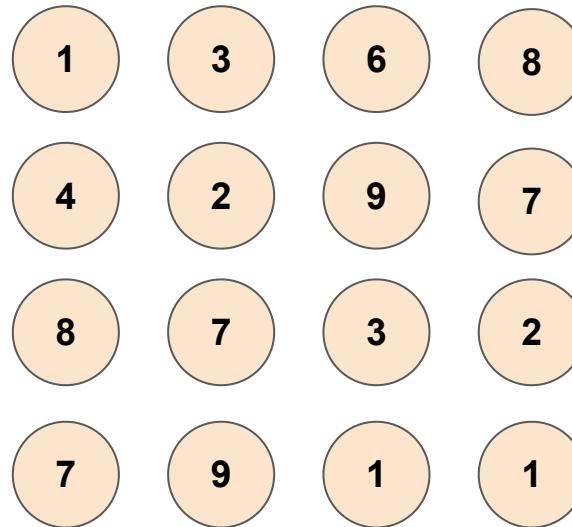
- Let's imagine a filter





# Deep Learning

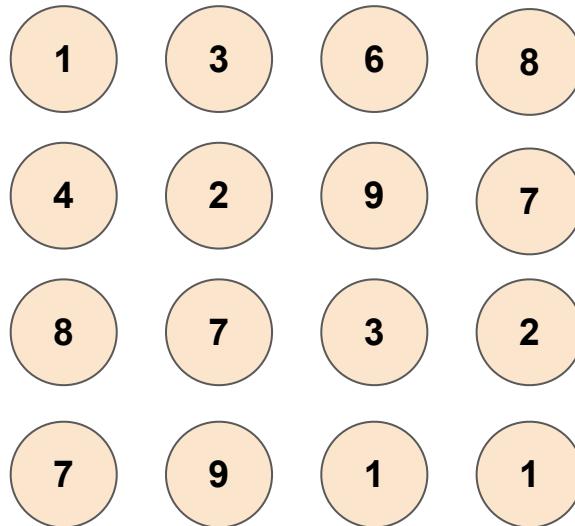
- We can reduce the size with subsampling





# Deep Learning

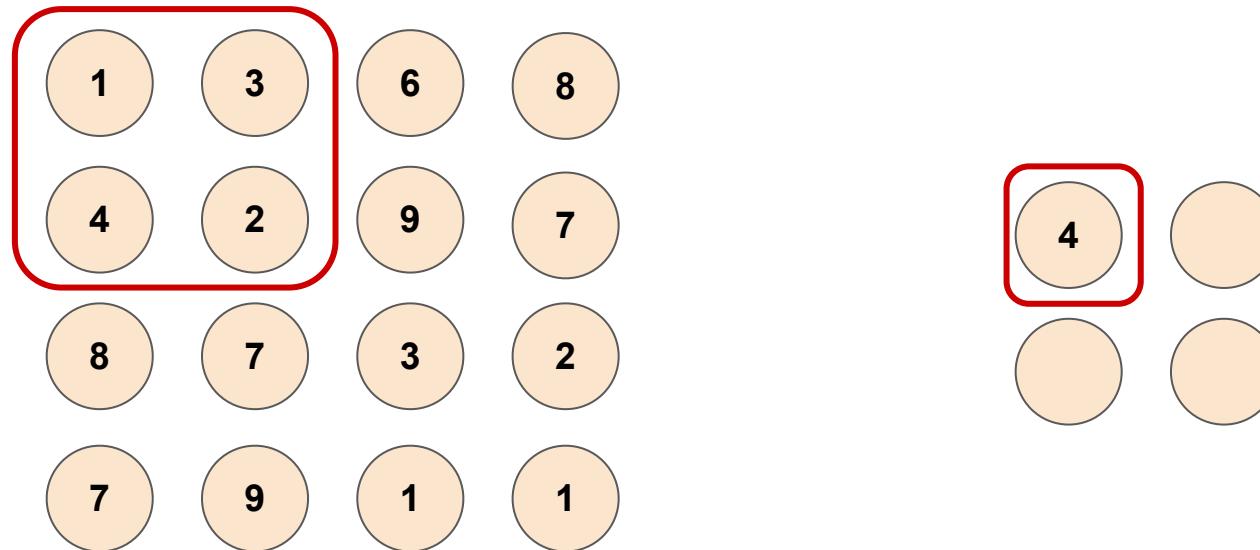
- We can reduce the size with subsampling





# Deep Learning

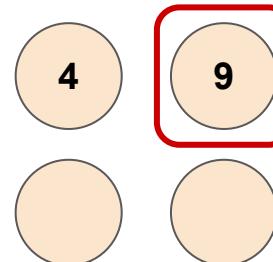
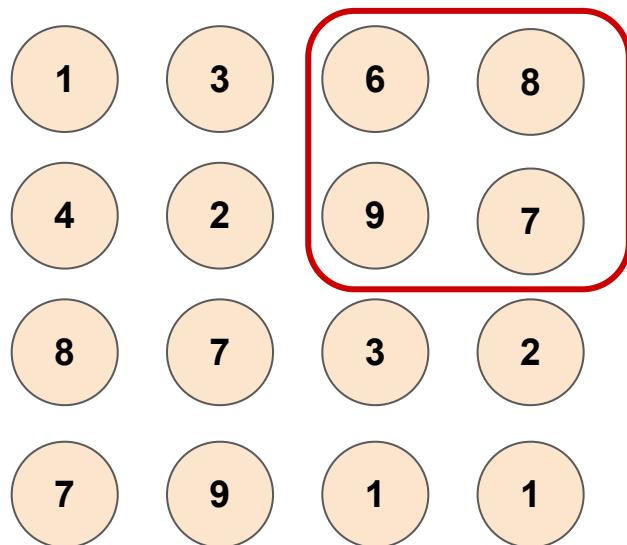
- Max Pooling: Window 2x2 , Stride: 2





# Deep Learning

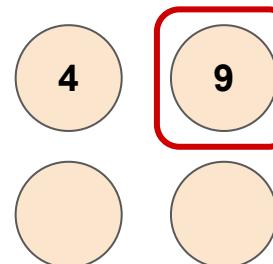
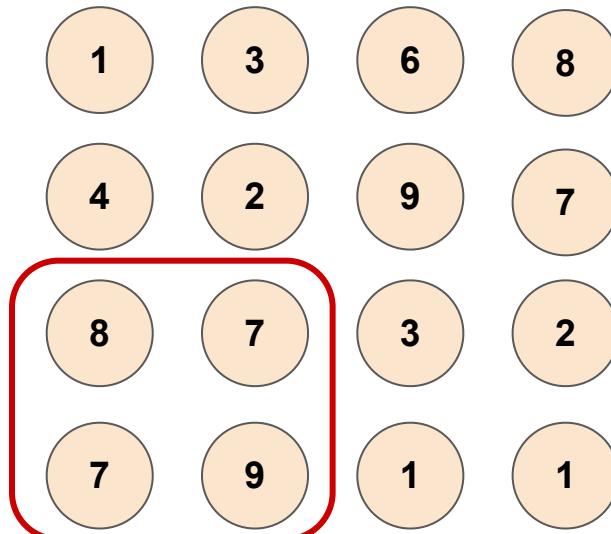
- Max Pooling: Window 2x2 , Stride: 2





# Deep Learning

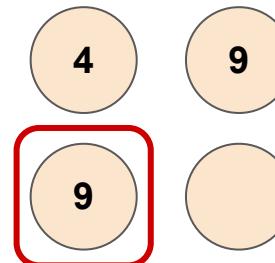
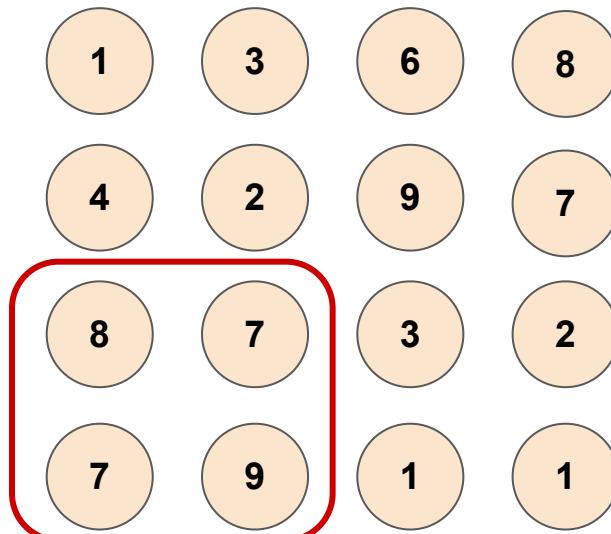
- Max Pooling: Window 2x2 , Stride: 2





# Deep Learning

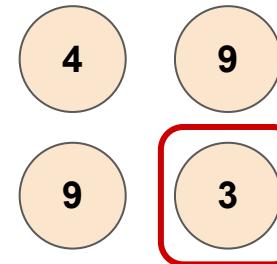
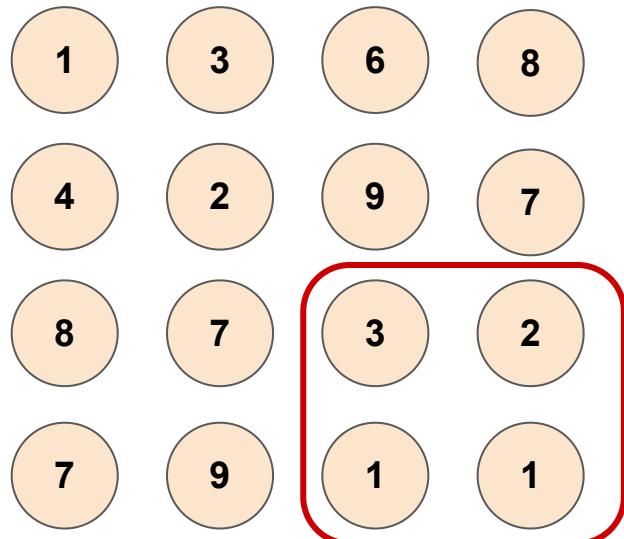
- Max Pooling: Window 2x2 , Stride: 2





# Deep Learning

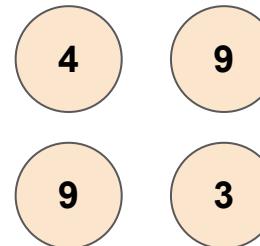
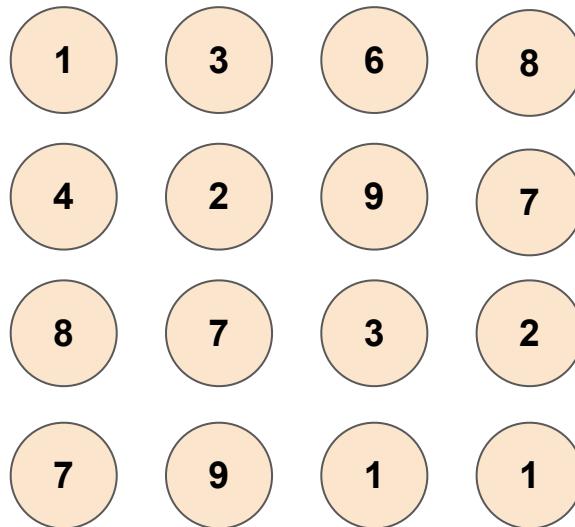
- Max Pooling: Window 2x2 , Stride: 2





# Deep Learning

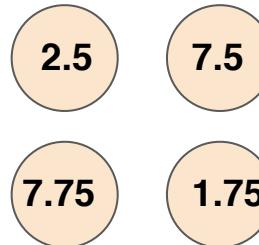
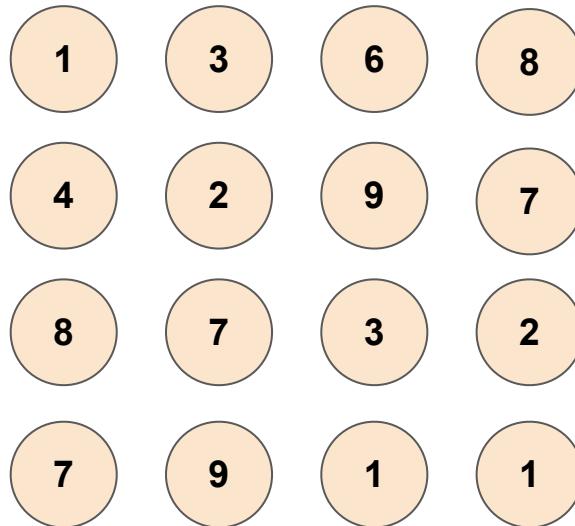
- Max Pooling: Window 2x2 , Stride: 2





# Deep Learning

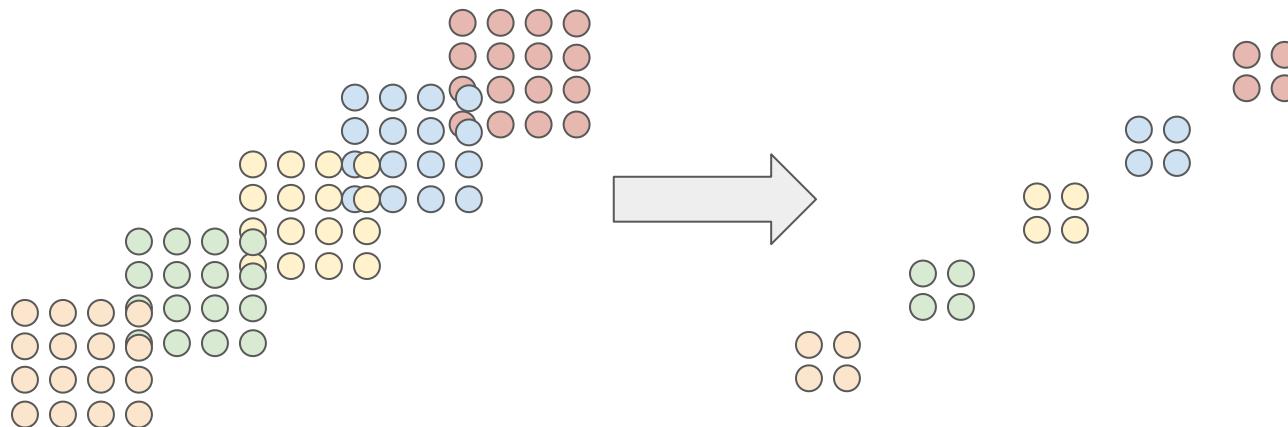
- Average Pooling: Window 2x2 , Stride: 2





# Deep Learning

- This greatly reduces our number of parameters!





# Deep Learning

- This pooling layer will end up removing a lot of information, even a small pooling “kernel” of 2 by 2 with a stride of 2 will remove 75% of the input data.



# Deep Learning

- Another common technique deployed with CNN is called “Dropout”
- Dropout can be thought of as a form of regularization to help prevent overfitting.



# Deep Learning

- This helps prevent units from “co-adapting” too much.
- Let’s also quickly point out some famous CNN architectures!



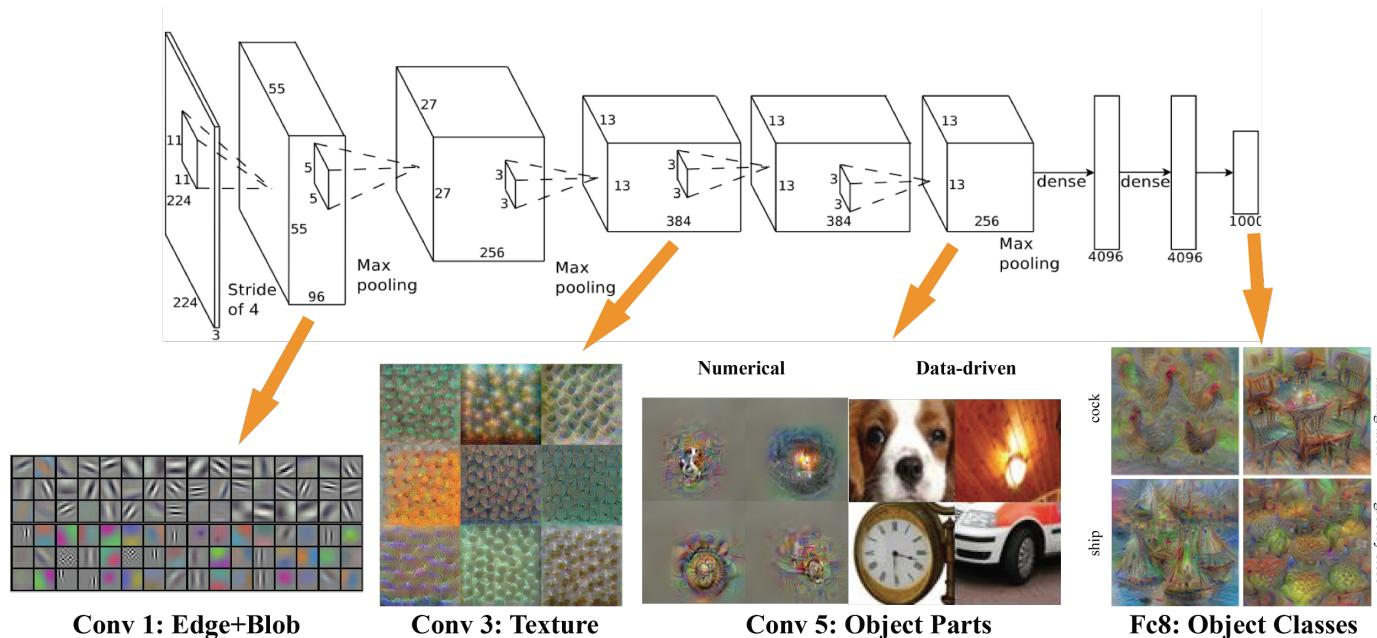
# Deep Learning

- LeNet-5 by Yann LeCun
- AlexNet by Alex Krizhevsky et al.



# Deep Learning

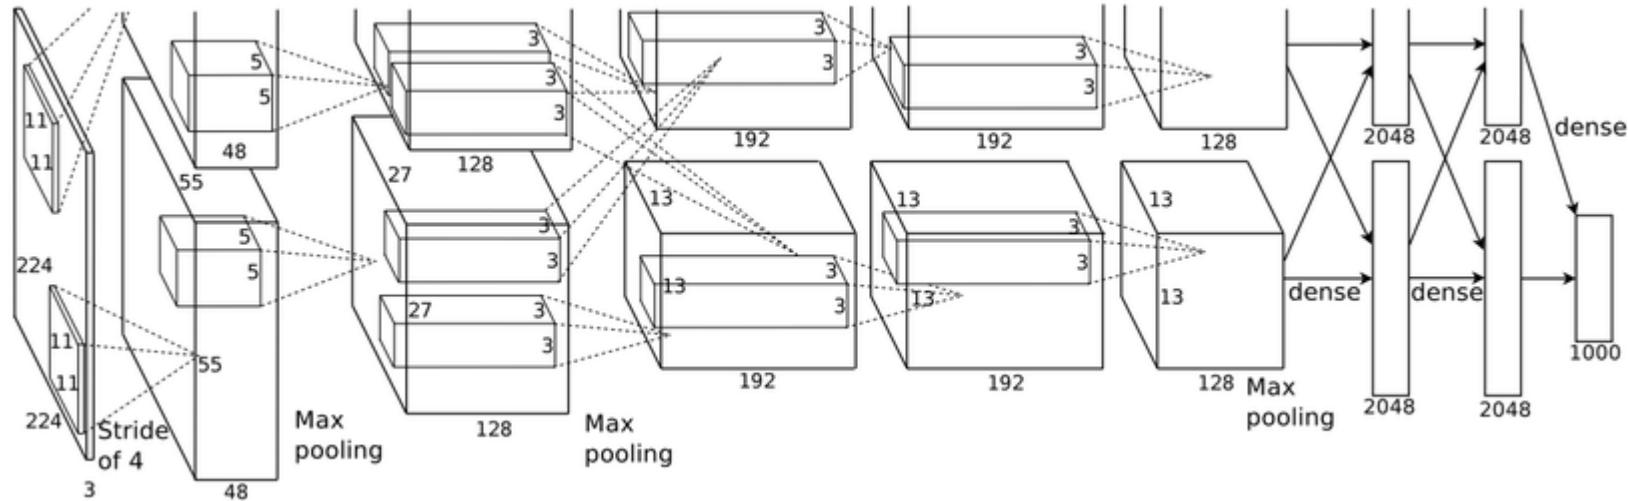
- AlexNet





# Deep Learning

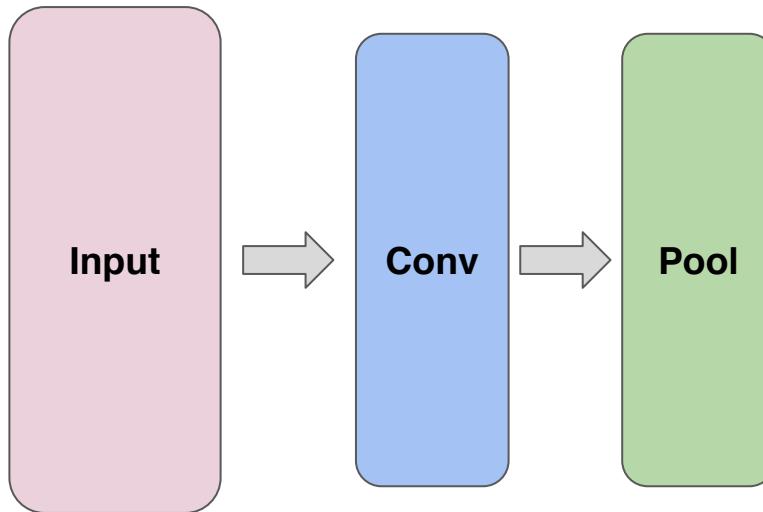
- AlexNet





# Deep Learning

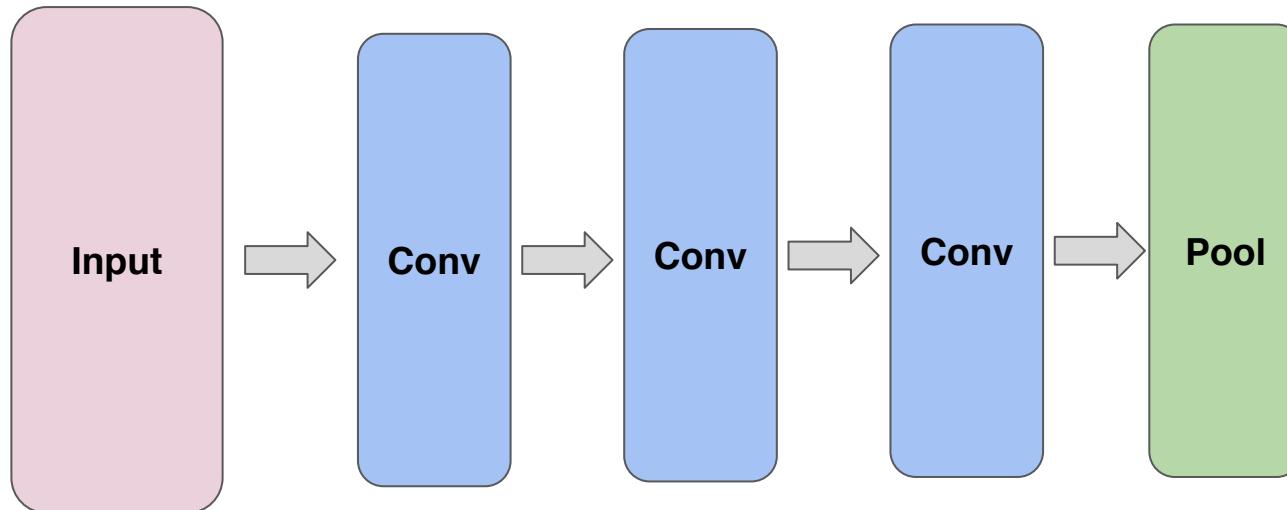
- CNNs can have all types of architectures!





# Deep Learning

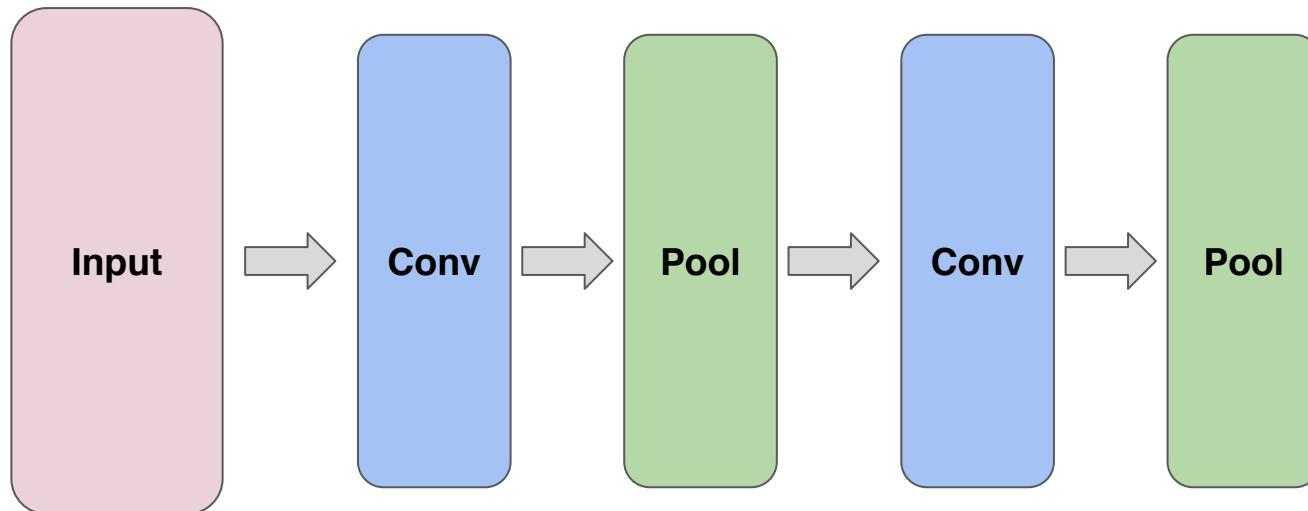
- CNNs can have all types of architectures!





# Deep Learning

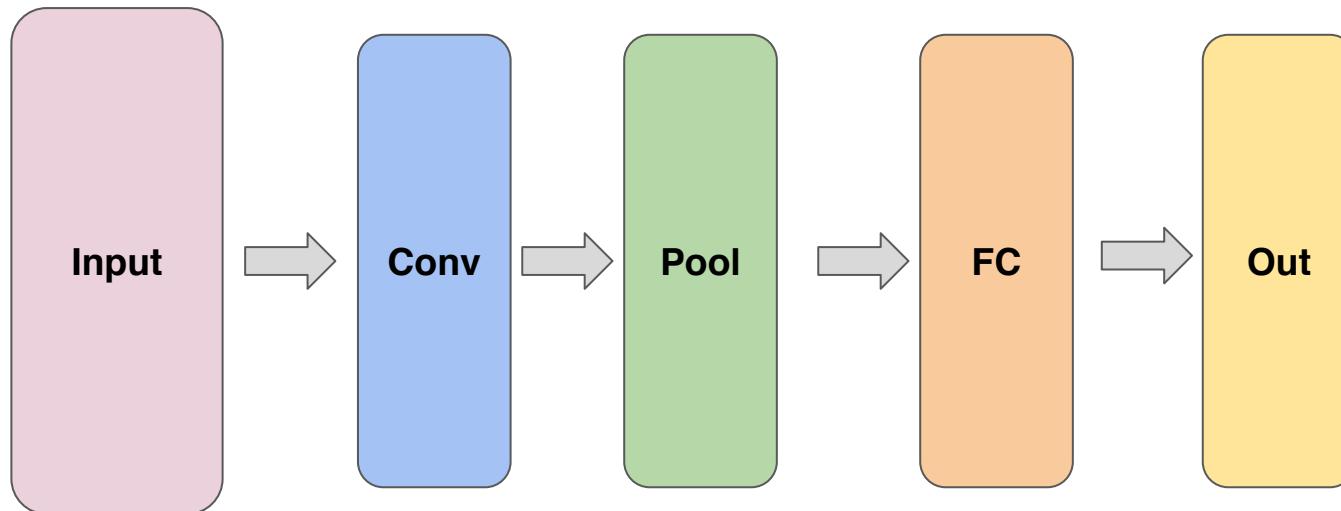
- CNNs can have all types of architectures!





# Deep Learning

- CNNs can have all types of architectures!





# MNIST Data Revisited



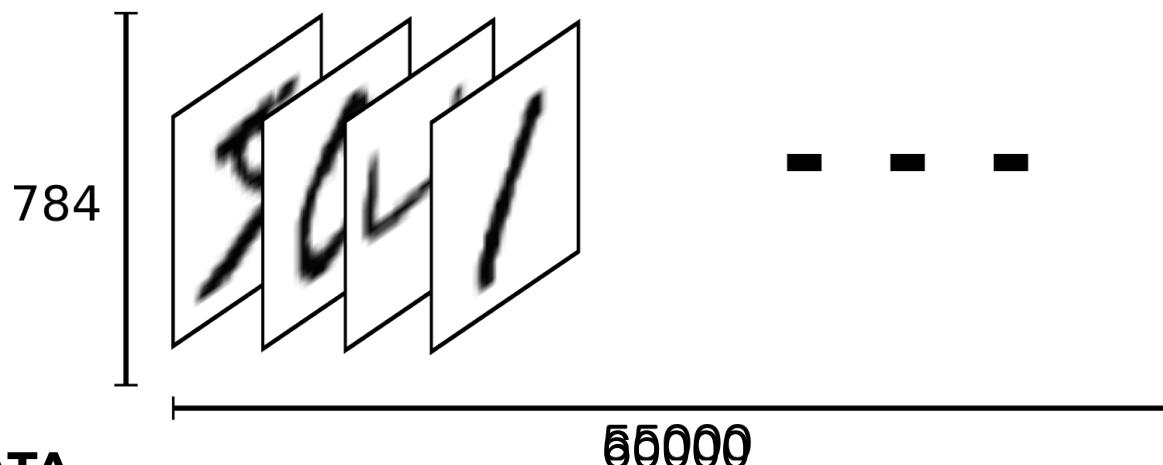
# Deep Learning

- Recall that flattening out the MNIST data caused us to lose 2D information.
- With CNNs, we can feed in the data as an array of 2D images.



## Deep Learning

- We can think of the entire group of the 60,000 images as a tensor (an n-dimensional array)





# Deep Learning

- For the labels we'll use One-Hot Encoding.
- This means that instead of having labels such as “One”, “Two”, etc... we'll have a single array for each image.



# Deep Learning

- The label is represented based off the index position in the label array.
- The corresponding label will be a 1 at the index location and zero everywhere else.
- For example, 4 would have this label array:
  - [0,0,0,0,1,0,0,0,0,0]



# Deep Learning

- Keep in mind that when dealing with tensors of image data, we actually end up with 4 dimensions:
  - Number of Images
  - Height
  - Width
  - Color Channels



# MNIST with CNN

PART ONE - The Data



# MNIST with CNN

PART TWO - Model and Training



# MNIST with CNN

PART THREE - Model Evaluation



# CIFAR-10 with CNN

PART ONE - Data



# Deep Learning

- CIFAR-10 dataset are 32by32 images of 10 different objects:
  - Airplane,Car,Bird,Cat,Deer,Dog,Frog, Horse,Ship,Truck
  - These are color images!



# Deep Learning

- In this series of lectures we will be reusing a lot of the previous CNN code, so we will only focus on new additions due to the introduction of 3 color channels (RGB).



# CIFAR-10 with CNN

PART TWO - Model Evaluation



# Downloading Image Data



# Deep Learning

- So far we've only dealt with pre-packaged data sets such as MNIST and CIFAR-10.
- But what about working with real image files?  
Like .jpg or .png?



# Deep Learning

- Let's learn about TensorFlow's built in tools for generating image data batches from directories with real files.
- First we need to download a large zip file of data.



# Deep Learning

- This dataset is available as a zip file called **cell\_images.zip**
- The link is a supplemental resource in this lecture, it is also within the same Google Drive resource as the slides.
- Let's explore getting the data!



# CNN on Custom Images

## Part 1 - The Data



# **CNN on Custom Images**

## **Part 2 ImageDataGenerator**



# CNN on Custom Images Part 3

## Creating the Model



# CNN on Custom Images Part 4

## Model Evaluation



# Let's get started!



# CNN Exercises



# CNN Exercises Solutions